

```
In [21]: import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn import metrics

# Load the dataset
crops = pd.read_csv("data/soil_measures.csv")
```

Check for missing values.

```
In [3]: crops.isna().sum()
```

```
Out[3]: N      0
P      0
K      0
ph     0
crop   0
dtype: int64
```

```
In [4]: crops.dtypes
```

```
Out[4]: N      int64
P      int64
K      int64
ph     float64
crop   object
dtype: object
```

```
In [5]: crops["crop"].value_counts()
```

```
Out[5]: crop
        rice      100
        maize     100
        chickpea   100
        kidneybeans 100
        pigeonpeas 100
        mothbeans  100
        mungbean    100
        blackgram   100
        lentil      100
        pomegranate 100
        banana      100
        mango       100
        grapes      100
        watermelon  100
        muskmelon   100
        apple       100
        orange      100
        papaya      100
        coconut     100
        cotton      100
        jute        100
        coffee      100
        Name: count, dtype: int64
```

```
In [6]: features = ["N", "P", "K", "ph"]
```

```
In [7]: X = crops[features]
```

```
In [8]: y = crops[["crop"]]
```

```
In [9]: X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3)
```

Trying out logistic regression.

```
In [10]: log_reg = LogisticRegression()
```

```
In [ ]: log_reg.fit(X_train, y_train)
```

```
In [12]: y_pred = log_reg.predict(X_test)
```

```
In [13]: display(metrics.accuracy_score(y_test,y_pred))
```

0.5484848484848485

Not the best results. Now lets see how a random forest performs.

```
In [14]: forest = RandomForestClassifier()
```

```
In [ ]: forest.fit(X_train,y_train)
```

```
In [16]: y_pred = forest.predict(X_test)
display(metrics.accuracy_score(y_test,y_pred))
```

0.7636363636363637

Much better!

Let's find the most relevant feature.

```
In [ ]: feature_scores = {}
for feature in features:
    new_X=X[feature].values.reshape(-1,1)
    X_train_new, X_test_new, y_train_new, y_test_new = train_test_split(new_X, y)
    new_forest = RandomForestClassifier()
    new_forest.fit(X_train_new, y_train_new)
    y_pred_new = new_forest.predict(X_test_new)
    score = metrics.accuracy_score(y_test_new, y_pred_new)
    feature_scores[feature] = score
```

```
In [18]: print(feature_scores)
```

```
{'N': 0.12, 'P': 0.19090909090909092, 'K': 0.3054545454545455, 'ph': 0.13272727272727272}
```

On its own, the content ratio of potassium is the best indicator for the optimal crop choice.

Let's find the best set of parameters.

```
In [22]: param_grid = {
    'bootstrap': [True],
    'max_depth': [80, 90, 100, 110],
    'max_features': [2, 3],
    'min_samples_leaf': [3, 4, 5],
    'min_samples_split': [8, 10, 12],
    'n_estimators': [100, 200, 300, 1000]
}

forest = RandomForestClassifier()
grid = GridSearchCV(estimator=forest, param_grid=param_grid, cv=3, n_jobs=-1)
```

```
In [ ]: grid.fit(X_train, y_train)
```

```
In [28]: display(grid.best_params_)
```

```
{'bootstrap': True,
 'max_depth': 80,
 'max_features': 2,
 'min_samples_leaf': 3,
 'min_samples_split': 12,
 'n_estimators': 100}
```

Now that we have our best params, we select the estimator with the best results.

```
In [25]: best_forest = grid.best_estimator_
```

```
In [ ]: best_forest.fit(X_train,y_train)
```

```
In [27]: y_pred = best_forest.predict(X_test)
accuracy = metrics.accuracy_score(y_test, y_pred)

display(accuracy)
```

0.7757575757575758

A tiny improvement but an improvement nonetheless.